


# An example deep convolution network

- Input: 28x28 grayscale image
- Output: 10 classes. One output per class.

# A Convolutional Net

- Let's assume we have 28x28 grayscale images as input to our conv net. So we will input 28x28x1 samples into the net.
- Let's fix our kernel size at 5x5 and, to make this simple, pad our images with zeros and use a stride = 1.
- Let's use max pooling on the output, with a 2x2 pooling region and a stride of 2.
- Let's extract 32 features after the first layer.
- So the output from this layer will be 14x14x32.


# A Convolutional Net



# A Convolutional Net

- Now let's make a second layer, also convolutional.
- Let's fix our kernel size at  $5 \times 5$ , pad our images with zeros and use a stride = 1.
- Let's use max pooling on the output again, with a  $2 \times 2$  pooling region and a stride of 2.
- Let's extract 64 features after the second layer.
- So the output from this layer will be  $7 \times 7 \times 64$ .


# A Convolutional Net



# A Convolutional Net

- Our third layer will be a fully connected layer mapping our convolutional features to a 1024 dimensional feature space.
- This layer is just like any of the hidden layers you've seen before. It is a linear transformation followed by ReLU.
- So the output from this layer will be  $1 \times 1 \times 1024$ .

# A Convolutional Net




# A Convolutional Net

- Finally, we'll map this feature space to a 10 class output space and use a softmax with a MLE/cross entropy loss function.
- And...we're done!




# A Convolutional Net



$$\text{Parameters} = (5 \times 5 \times 1 \times 32 + 32) + (5 \times 5 \times 32 \times 64 + 64) + (7 \times 7 \times 64 \times 1024 + 1024) + (1024 \times 10 + 10)$$

# Convolutional Layer: Padding + Stride



$$\text{Output dimension} = (\text{input dim} - \text{kernel size} + 2 * \text{padding}) / \text{stride} + 1$$